

# MNIST Neural Network Implementation

Josue Rolando Naranjo Sieiro

*Student of Computer Science, University of Havana, Cuba*

November 23, 2024

## Abstract

This report summarizes the implementation and results of neural networks with varying layers to classify handwritten digits from the MNIST dataset. The experiments include a single-layer neural network, a two-layer neural network, and a three-layer neural network.

## 1 Introduction

The MNIST dataset is a well-known dataset for handwritten digit classification. This project implements three neural network architectures to classify the digits: a single-layer neural network, a two-layer neural network, and a three-layer neural network. The goal is to compare the performance and training time of these architectures.

## 2 Single Layer Neural Network

The single-layer neural network consists of one hidden layer and an output layer. The implementation is provided in the `NN1Layer.py` script. The following functions are used:

### 2.1 Functions

- `one_hot_encode(y)`: Converts labels to one-hot encoded vectors.
- `init_weights(input_size, hidden_size, output_size)`: Initializes weights and biases.
- `sigmoid(z)`: Sigmoid activation function.
- `sigmoid_derivative(a)`: Derivative of the sigmoid function.
- `relu(z)`: ReLU activation function.
- `relu_derivative(z)`: Derivative of the ReLU function.

- **softmax(z)**: Softmax activation function.
- **forward\_propagation(X, W1, b1, W2, b2)**: Performs forward propagation.
- **compute\_cost(A2, Y)**: Computes the cost using categorical cross-entropy.
- **backward\_propagation(X, Y, cache, W2)**: Performs backward propagation.
- **update\_parameters(W1, b1, W2, b2, gradients, learning\_rate)**: Updates weights and biases using gradient descent.
- **train(X, Y, input\_size, hidden\_size, output\_size, epochs, learning\_rate)**: Trains the neural network.
- **predict(X, parameters)**: Makes predictions using the trained model.

## 2.2 Results

The single-layer neural network achieved an accuracy of 92.28% on the test set.

## 3 Two Layer Neural Network

The two-layer neural network consists of two hidden layers and an output layer. The implementation is provided in the `NN2Layer.py` script. The following functions are used:

### 3.1 Functions

- **one\_hot\_encode(y)**: Converts labels to one-hot encoded vectors.
- **init\_weights(layer\_sizes)**: Initializes weights and biases.
- **relu(z)**: ReLU activation function.
- **relu\_derivative(z)**: Derivative of the ReLU function.
- **softmax(z)**: Softmax activation function.
- **forward\_propagation(X, parameters)**: Performs forward propagation.
- **compute\_cost(AL, Y, parameters, lambd)**: Computes the cost with L2 regularization.
- **backward\_propagation(Y, cache, parameters, lambd)**: Performs backward propagation.
- **update\_parameters(parameters, gradients, optimizer\_params)**: Updates weights and biases using the Adam optimizer.

- **initialize\_optimizer\_params(parameters)**: Initializes parameters for the Adam optimizer.
- **get\_mini\_batches(X, Y, batch\_size)**: Generates mini-batches for training.
- **train(X, Y, layer\_sizes, epochs, lambd, batch\_size)**: Trains the neural network.
- **predict(X, parameters)**: Makes predictions using the trained model.

### 3.2 Results

The two-layer neural network achieved an accuracy of 97.71% on the test set.

## 4 Three Layer Neural Network

The three-layer neural network consists of three hidden layers and an output layer. The implementation is provided in the `NN3Layer.py` script. The following functions are used:

### 4.1 Functions

- **one\_hot\_encode(y)**: Converts labels to one-hot encoded vectors.
- **init\_weights(layer\_sizes)**: Initializes weights and biases.
- **relu(z)**: ReLU activation function.
- **relu\_derivative(z)**: Derivative of the ReLU function.
- **softmax(z)**: Softmax activation function.
- **forward\_propagation(X, parameters)**: Performs forward propagation.
- **compute\_cost(AL, Y, parameters, lambd)**: Computes the cost with L2 regularization.
- **backward\_propagation(Y, cache, parameters, lambd)**: Performs backward propagation.
- **update\_parameters(parameters, gradients, optimizer\_params)**: Updates weights and biases using the Adam optimizer.
- **initialize\_optimizer\_params(parameters)**: Initializes parameters for the Adam optimizer.
- **get\_mini\_batches(X, Y, batch\_size)**: Generates mini-batches for training.

- **train(X, Y, layer\_sizes, epochs, lambd, batch\_size)**: Trains the neural network.
- **predict(X, parameters)**: Makes predictions using the trained model.

## 4.2 Results

The three-layer neural network achieved an accuracy of 97.70% on the test set, but with a significantly longer training time compared to the two-layer neural network.

## 5 Conclusion

The experiments demonstrate that increasing the number of layers in a neural network can improve accuracy, but it also increases the training time. The two-layer neural network provided the best balance between accuracy and training time, achieving an accuracy of 97.71%.

## 6 References

- MNIST dataset: <http://yann.lecun.com/exdb/mnist/>
- Neural Network implementation: `NN1Layer.py`, `NN2Layer.py`, `NN3Layer.py`
- Pre-trained model: `trained_model.pkl`